
FIDELITY IS NOT UTILITY: A CONTROLLED DISSOCIATION IN LEARNED IMAGE TRANSLATION

A PREPRINT

Ishrith Gowda

University of California, Berkeley
ishrithgowda@berkeley.edu

Chunwei Liu

Purdue University

ABSTRACT

Image translation for harmonization and domain adaptation is routinely selected on distributional fidelity metrics such as the Fréchet Inception Distance (FID), on the assumption that a more target-like image is more useful: a team can lower FID and ship a model that is worse at the task. That fidelity can mis-rank generators is known; we make a sharper, mechanistic claim: the utility loss is a property of the learned distribution-matching *objective* itself, not of any architecture, which is why it survives the move from GANs to diffusion. Our instrument is a *frozen* downstream evaluator scored alongside FID on identical inputs, plus a matched non-learned control. On GTA5→Cityscapes ($N=1000$, frozen SegFormer), CycleGAN more than halves FID (137.2→65.5) while mIoU falls 0.380→0.167; the non-learned control preserves mIoU (0.363) without improving FID. Sweeping SDEdit’s noise strength traces a monotonic frontier: FID 145.1→88.0 as mIoU falls 0.323→0.106, with seed standard deviation at most 0.005 mIoU. Two controls close the obvious escapes: with an empty prompt, fidelity is worse than the untranslated source (FID 154.2) yet mIoU still collapses (0.170), decoupling the damage from any fidelity gain; and edge-conditioned ControlNet does not escape the frontier (mIoU 0.221 at FID 133.2, versus 0.289 at matched FID). Multi-site brain-tumor MRI reproduces it: frozen nnU-Net tumor Dice falls 0.682→0.064 under learned harmonization but holds at 0.669 non-learned. Across both generative families and both domains, the low-FID/high-utility corner stayed empty. We do not claim fidelity is uninformative, only that it is non-monotone in utility once a learned prior enters the loop; our contribution is a measurement and a protocol, not a method.

1 Introduction

Distribution-matching translation is trained and evaluated to make outputs *look like* a target domain. The implicit promise is that looking more like the target makes outputs more *useful*—for example, that harmonizing site B to site A helps a model trained on A . We test that promise directly and find it false for learned translation.

Our contributions are:

- A **controlled protocol** for separating fidelity from utility: a frozen downstream evaluator plus a matched non-learned control, measuring FID and the task metric on identical inputs (Section 3).
- Evidence that the fidelity–utility dissociation is a property of the **learned distribution-matching objective**, not a single architecture—it reproduces under both a GAN (CycleGAN) and a diffusion model (SDEdit).
- A **continuous characterization** via SDEdit strength: a monotonic FID-vs-mIoU frontier with tight seed variance, an empty-prompt ablation that decouples utility loss from fidelity gain, and a structure-preserving (ControlNet) baseline that fails to escape the frontier.
- **Cross-domain** generality: the same dissociation in driving scenes and in clinical brain-tumor MRI.

2 Related Work

Three separately-established observations converge in our claim; we unify them.

Fidelity metrics mis-rank generators for downstream tasks. Konz et al. [2024] show no perceptual metric reliably tracks segmentation quality in medical translation; Dutt et al. [2025] show low FID does not predict clinical utility; and Kynkäänniemi et al. [2023] show FID is gameable by aligning ImageNet-class statistics. These are *metric-evaluation* studies; we add a learned/non-learned control and a second generative family and domain.

Distribution-matching translation alters task content. Cohen et al. [2018] show CycleGAN-style matching can hallucinate or remove tumors to fit the target distribution, and semantic-consistency losses [Hoffman et al., 2018, Peng et al., 2023] exist precisely to patch this. We show that *without* such non-distributional anchoring the pure learned objective degrades utility, reframing those losses as evidence for our thesis.

Non-learned alignment rivals learned translation. Yang and Soatto [2020] reach strong segmentation adaptation with a non-learned spectral swap. We use histogram/color transfer as the analogous control and, unlike FDA, measure the fidelity axis explicitly to expose the dissociation. Richter et al. [2022] likewise treat a segmentation-stability metric as a guardrail that realism must not break.

Positioning. No prior work runs the clean 2×2 {learned, non-learned} $\times$ {fidelity $\downarrow$?, utility $\downarrow$?} as its thesis, ties the failure to the objective across both a GAN *and* a diffusion model, and does so across driving and medicine. This is a measurement/positioning contribution, not a new architecture.

3 Setup: Separating Fidelity from Utility

Given source images x and a translation T , we evaluate $T(x)$ on two axes. **Fidelity** is $\text{FID}(T(x), \text{target})$ via clean-fid [Parmar et al., 2022]; lower is closer to the target. **Utility** is a *frozen* downstream model f (never retrained) applied to $T(x)$ and scored against ground truth; because f is fixed, any change in its score is attributable purely to T . We compare four families of T : (i) identity (raw source, reference); (ii) non-learned (histogram / color transfer); (iii) learned-GAN (CycleGAN, Zhu et al., 2017); and (iv) learned-diffusion (SDEdit, Meng et al., 2022, over a sweep of noise strengths, plus ControlNet-Canny, Zhang et al., 2023). FID is compared only at equal sample size, which it is known to bias.

4 Experiments

4.1 Driving: GTA5 $\rightarrow$ Cityscapes

The source is GTA5 [Richter et al., 2016] and the target is Cityscapes [Cordts et al., 2016]. The frozen evaluator is a Cityscapes-finetuned SegFormer-b4 [Xie et al., 2021] (19 classes); utility is mIoU. We use $N=1000$ GTA5 images per condition, with the Cityscapes validation set (500 images) as the fixed FID reference, and identical evaluator and clean-fid (GPU) for all conditions. SDEdit runs on Stable Diffusion [Rombach et al., 2022] and rows report mean $\pm$ std over three seeds {42, 43, 44}. Results are in Table 1 and Figure 1.

Findings. Figure 2 shows representative translations. (a) Learned translation lowers FID relative to raw while collapsing mIoU by 15–72%; CycleGAN more than halves FID (137 $\rightarrow$ 65) at mIoU 0.167. (b) The non-learned baseline does not improve FID yet preserves mIoU (−1.8%). (c) SDEdit strength traces a **monotonic frontier** with tight seed variance ($\sigma(\text{mIoU}) \leq 0.005$): as strength rises, FID falls (145 $\rightarrow$ 88) and mIoU falls with it (0.323 $\rightarrow$ 0.106); no learned setting reaches the ideal corner, occupied only by raw/non-learned. (d) **Empty-prompt ablation.** With no text prompt, FID is *worse* than raw (154.2 $>$ 137.2) yet mIoU still collapses (0.170)—the utility loss is the diffusion prior itself, not text guidance, and occurs even absent any fidelity gain. (e) **Structure-preserving baseline.** Conditioning on the source image’s Canny edges (ControlNet-Canny) does *not* rescue utility: it lands on the learned frontier (FID 133.2, mIoU 0.221 $\pm$ 0.002), Pareto-dominated by SDEdit at equal FID (SDEdit 0.30: FID 133.1, mIoU 0.289). Edge structure is not the task-relevant signal; only conditioning on the segmentation *map* [Peng et al., 2023] recovers utility, and only by injecting the labels—the non-distributional constraint our thesis requires. A local $N=200$ run reproduces the same ordering on an independent machine.

Table 1: GTA5→Cityscapes, $N=1000$, frozen SegFormer-b4 + clean-fid. SDEdit rows are mean $\pm$ std over three seeds. Learned translation lowers FID while collapsing mIoU; the non-learned baseline preserves mIoU without improving FID; no learned setting reaches the low-FID/high-mIoU corner.

condition	family	FID $\downarrow$	mIoU $\uparrow$	Δ mIoU
raw GTA5	baseline	137.2	0.380	—
color match	non-learned	165.9	0.363	−0.018
SDEdit 0.20	diffusion	145.1 $\pm$ 1.8	0.323 $\pm$ 0.003	−0.057
SDEdit 0.30	diffusion	133.1 $\pm$ 2.1	0.289 $\pm$ 0.002	−0.092
SDEdit 0.40	diffusion	116.4 $\pm$ 1.1	0.237 $\pm$ 0.003	−0.143
SDEdit 0.50	diffusion	100.3 $\pm$ 1.4	0.194 $\pm$ 0.005	−0.187
SDEdit 0.60	diffusion	90.8 $\pm$ 2.0	0.150 $\pm$ 0.001	−0.231
SDEdit 0.70	diffusion	88.0 $\pm$ 0.4	0.106 $\pm$ 0.001	−0.274
CycleGAN	learned-GAN	65.5	0.167	−0.214
SDEdit 0.50, empty prompt	diffusion (abl.)	154.2 $\pm$ 1.9	0.170 $\pm$ 0.003	−0.211
ControlNet-Canny	diffusion + edge	133.2 $\pm$ 0.3	0.221 $\pm$ 0.002	−0.159




Figure 1: The fidelity–utility plane for GTA5→Cityscapes (mean $\pm$ std over seeds). The SDEdit strength sweep traces a monotonic frontier: as strength rises, FID falls and mIoU falls with it. Every learned point (CycleGAN, all SDEdit strengths, the empty-prompt ablation, and edge-conditioned ControlNet) sits well below raw and the non-learned baseline on utility; the ideal low-FID/high-mIoU corner is empty.

4.2 Medical: Multi-Site Brain-Tumor MRI

The frozen evaluator is an nnU-Net [Isensee et al., 2021] trained on BraTS [Menze et al., 2015] and frozen; utility is tumor Dice on the external UPenn-GBM site [Bakas et al., 2022] ($n=103$). The learned harmonizer is SA-CycleGAN-2.5D, a self-attention CycleGAN over 2.5D multi-modal MRI. Table 2 reports the result.

Findings. Learned harmonization collapses external-site Dice ($0.682 \rightarrow 0.064$) while the non-learned baseline preserves it (0.669). Content/segmentation-consistency losses reduce but do not eliminate the loss ($0.205 \ll 0.682$), consistent with the driving result that the degradation is intrinsic to the learned objective. A second external site (RHUH-GBM, Cepeda et al., 2023, $n=40$) confirms the direction ($0.718 \rightarrow 0.599$, -16.6%). We note two evaluation




Figure 2: Qualitative GTA5→Cityscapes translations (three scenes, seed 42). Left to right: raw source, non-learned color match, CycleGAN, SDEdit at increasing strength, the empty-prompt SDEdit, and edge-conditioned ControlNet. As the learned translations become more Cityscapes-like they visibly hallucinate and alter scene structure—buildings, lane markings, added objects—that the frozen segmenter depends on, while the non-learned baseline leaves structure intact.

Table 2: External UPenn-GBM tumor Dice under a frozen BraTS-trained nnU-Net. Learned harmonization collapses Dice even though it makes images more source-like; the non-learned baseline preserves it.

input to frozen nnU-Net	family	Dice (fg. mean)
BraTS test (in-domain ref.)	—	0.771
UPenn raw (external)	baseline	0.682
UPenn + histogram match	non-learned	0.669
UPenn + learned harmonization	learned-GAN	0.064
UPenn + symmetric seg-consistency	learned-GAN (+ content)	0.205

protocols exist: the headline numbers use the frozen independent nnU-Net; a separate method-development evaluator ranks the symmetric loss as the best-preserving *learned* harmonizer, which belongs to a companion method paper.

4.3 Where the damage falls

The headline mIoU collapse hides structure. Breaking IoU out per class and grouping the 19 Cityscapes classes into large amorphous regions and thin, small, label-carrying objects (Table 3) shows the degradation is highly *selective*. The non-learned control loses a few percent on both groups, almost uniformly. Every learned translation instead loses far more on the thin classes than on the large ones: SDEdit at strength 0.50 costs 27% of large-region IoU but 83% of thin-object IoU, and CycleGAN 50% against 70%. Poles, traffic lights, traffic signs, riders and motorcycles are precisely the structures a segmenter needs and precisely what the learned prior does not preserve. Edge-conditioned ControlNet sits in between (27% / 63%): conditioning on edges recovers some thin structure, consistent with Section 3’s argument that what is missing is a non-distributional constraint, but not enough to escape the frontier.

5 Analysis

The learned objective optimizes appearance under a distribution/adversarial or denoising target with no term guaranteeing task-relevant signal survives. SDEdit strength is precisely the amount of learned prior injected, and utility falls monotonically with it. The empty-prompt result decouples utility loss from fidelity gain: passing through the prior damages task content even when the output is *less* target-like. Non-learned transforms, lacking a learned prior, cannot introduce this failure; edge-conditioned ControlNet, injecting structure but not semantics, does not avoid it.

Table 3: Per-class IoU grouped into large amorphous regions (road, sidewalk, building, vegetation, terrain, sky) and thin/small label-carrying objects (pole, traffic light, traffic sign, rider, motorcycle); classes whose raw IoU is below 0.05 are excluded, since a ratio off a near-zero baseline is noise. The non-learned control degrades both groups mildly and almost uniformly; every learned translation degrades the thin classes far more than the large ones.

group	raw	color match	CycleGAN	SDEdit 0.50	ControlNet
large regions	0.661	0.630 (−5%)	0.328 (−50%)	0.481 (−27%)	0.482 (−27%)
thin / small objects	0.174	0.168 (−4%)	0.052 (−70%)	0.030 (−83%)	0.065 (−63%)

6 Limitations

FID is sample-size biased; all conditions use the same N (1000 translated vs the 500-image Cityscapes validation set as the fixed reference), and SDEdit points report mean $\pm$ std over three seeds ($\sigma(\text{mIoU}) \leq 0.005$, $\sigma(\text{FID}) \leq 2.1$), so the ordering is robust to both sample size and seed. The structure-preserving baseline uses edge (Canny) conditioning; a segmentation-map-conditioned variant would recover utility only by injecting the labels, which concedes the thesis rather than refuting it.

7 Conclusion

Fidelity is not utility. Rewarding distributional realism can actively harm the downstream task for learned translation, across architectures and domains. We recommend reporting a frozen-evaluator utility axis alongside any fidelity metric, and treating non-learned baselines as first-class controls.

Reproducibility Statement

All conditions regenerate from a single command; seeds are fixed and logged; frozen evaluators are referenced by checkpoint id (SegFormer-b4, nnU-Net); FID uses clean-fid at a fixed N . Code, configs, and the $N=1000$ result set are released with the paper.

A Full per-class IoU

Table 4 reports every Cityscapes class, the breakdown behind Table 3. SDEdit and ControlNet columns are means over the three seeds. Classes with very small raw IoU (*train*, *bicycle*) are near-absent in the GTA5 subset; a relative change computed off such a baseline is noise, which is why the grouped statistics in Table 3 exclude classes with raw IoU below 0.05. The selectivity is visible directly: *road* loses little under every condition, while *traffic light*, *traffic sign*, *person* and *rider* are almost entirely destroyed by the learned translations and almost untouched by the non-learned control.

References

- Spyridon Bakas et al. The university of pennsylvania glioblastoma (UPenn-GBM) cohort: Advanced MRI, clinical, genomics, and radiomics. *Scientific Data*, 9(453), 2022.
- Santiago Cepeda et al. RHUH-GBM: A radiomics and clinical dataset of glioblastoma from río hortega university hospital. *Data in Brief*, 2023.
- Joseph Paul Cohen, Margaux Luck, and Sina Honari. Distribution matching losses can hallucinate features in medical image translation. In *Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, 2018.
- Marius Cordts, Mohamed Omran, Sebastian Ramos, Timo Rehfeld, Markus Enzweiler, Rodrigo Benenson, Uwe Franke, Stefan Roth, and Bernt Schiele. The Cityscapes dataset for semantic urban scene understanding. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- Raman Dutt et al. CheXGenBench: A unified benchmark for fidelity, privacy and utility of synthetic chest radiographs. *arXiv preprint arXiv:2505.10496*, 2025.
- Judy Hoffman, Eric Tzeng, Taesung Park, Jun-Yan Zhu, Phillip Isola, Kate Saenko, Alexei Efros, and Trevor Darrell. CyCADA: Cycle-consistent adversarial domain adaptation. In *International Conference on Machine Learning (ICML)*, 2018.

Table 4: Per-class IoU under a frozen SegFormer-b4, $N=1000$.

class	raw	color match	CycleGAN	SDEdit 0.50	ControlNet
road	0.937	0.932	0.843	0.845	0.884
sidewalk	0.540	0.502	0.202	0.280	0.309
building	0.711	0.681	0.356	0.504	0.505
wall	0.356	0.315	0.089	0.077	0.094
fence	0.102	0.111	0.020	0.021	0.031
pole	0.269	0.252	0.112	0.095	0.159
traffic light	0.073	0.067	0.008	0.008	0.012
traffic sign	0.069	0.063	0.012	0.017	0.022
vegetation	0.649	0.625	0.364	0.455	0.427
terrain	0.227	0.159	0.083	0.050	0.011
sky	0.901	0.883	0.121	0.750	0.755
person	0.270	0.240	0.035	0.024	0.111
rider	0.182	0.156	0.028	0.015	0.085
car	0.737	0.725	0.384	0.226	0.447
truck	0.666	0.608	0.198	0.197	0.175
bus	0.248	0.259	0.206	0.097	0.127
train	0.002	0.002	0.000	0.001	0.001
motorcycle	0.279	0.303	0.101	0.016	0.046
bicycle	0.008	0.007	0.001	0.001	0.004

- Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods*, 18(2):203–211, 2021.
- Nicholas Konz, Yuwen Chen, Hanxue Gu, Haoyu Dong, and Maciej A. Mazurowski. Rethinking perceptual metrics for medical image translation. *arXiv preprint arXiv:2404.07318*, 2024.
- Tuomas Kynkäänniemi, Tero Karras, Miika Aittala, Timo Aila, and Jaakko Lehtinen. The role of ImageNet classes in Fréchet inception distance. In *International Conference on Learning Representations (ICLR)*, 2023.
- Chenlin Meng, Yutong He, Yang Song, Jiaming Song, Jiajun Wu, Jun-Yan Zhu, and Stefano Ermon. SDEdit: Guided image synthesis and editing with stochastic differential equations. In *International Conference on Learning Representations (ICLR)*, 2022.
- Bjoern H. Menze et al. The multimodal brain tumor image segmentation benchmark (BRATS). *IEEE Transactions on Medical Imaging*, 34(10):1993–2024, 2015.
- Gaurav Parmar, Richard Zhang, and Jun-Yan Zhu. On aliased resizing and surprising subtleties in GAN evaluation. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Duo Peng, Ping Hu, Qihong Ke, and Jun Liu. Diffusion-based image translation with label guidance for domain adaptive semantic segmentation. In *IEEE International Conference on Computer Vision (ICCV)*, 2023.
- Stephan R. Richter, Vibhav Vineet, Stefan Roth, and Vladlen Koltun. Playing for data: Ground truth from computer games. In *European Conference on Computer Vision (ECCV)*, 2016.
- Stephan R. Richter, Hassan Abu AlHaija, and Vladlen Koltun. Enhancing photorealism enhancement. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2022.
- Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- Enze Xie, Wenhai Wang, Zhiding Yu, Anima Anandkumar, Jose M. Alvarez, and Ping Luo. SegFormer: Simple and efficient design for semantic segmentation with transformers. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2021.
- Yanchao Yang and Stefano Soatto. FDA: Fourier domain adaptation for semantic segmentation. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *IEEE International Conference on Computer Vision (ICCV)*, 2023.

Jun-Yan Zhu, Taesung Park, Phillip Isola, and Alexei A. Efros. Unpaired image-to-image translation using cycle-consistent adversarial networks. In *IEEE International Conference on Computer Vision (ICCV)*, 2017.